

Python Programming and Basic Data Science

Shuvro Pal

Data Scientist, Markopolo.ai

AI Team Lead, Zantrik

ML Facilitator, Crowdsourcing by Google

Clustering

What is it?

- Clustering is a technique for finding similarity groups in data, called clusters.
- A good clustering method will produce clusters with
 - Inter-clusters distance → maximized
 - Intra-clusters distance → minimized

k-means

What is it?

- K-means clustering is an unsupervised learning algorithm used to group similar data points into clusters. It tries to partition the dataset into K clusters, where each data point belongs to the cluster with the nearest centroid.
- **Key idea:** The k-means algorithm aims to group data points that are similar to each other while ensuring that each group (cluster) is as distinct as possible from the others.

Example: Imagine you have a group of fruits with different sizes and colors. K-means clustering helps group similar fruits together based on these features (e.g., grouping apples, oranges, and bananas separately).

k-means

Use cases?

- Customer segmentation in marketing
- Image compression
- Document clustering and topic detection
- Grouping similar products or services

k-means

Real World Examples?

Problem: Images are made up of thousands of pixels, each with an RGB value. Storing or transmitting high-resolution images can be inefficient due to their size.

K-means Solution:

- K-means clustering is used to reduce the number of colors in the image. Each pixel is assigned to the nearest cluster (representing a color).
 - Instead of storing the exact RGB value for every pixel, we store only the cluster ID and the centroid (color) for that cluster, drastically reducing the image size.
1. **Before K-means:** A high-resolution image might have millions of unique colors (pixel values).
 2. **After K-means:** K-means reduces the number of colors to a smaller set of K colors (e.g., K=16). The resulting image looks similar but takes up far less storage space.

Impact: K-means allows for efficient image compression without significant loss in visual quality. This is used in JPEG compression and web image optimization.

k-means

How it actually works?

Step 1: Choose the number of clusters (K). Decide how many clusters you want to divide your data into (this is a key input to the algorithm).

Step 2: Initialize centroids. Randomly assign K points as the starting centroids.

Step 3: Assign data points to clusters. Each data point is assigned to the nearest centroid.

Step 4: Update centroids. Calculate the new centroids by averaging the data points in each cluster.

Step 5: Repeat steps 3 and 4 until the centroids no longer move significantly (convergence).

k-means

Some math please?

Use the k-means algorithm and Euclidean distance to cluster the following 8 examples into 3($k=3$) clusters:

$A_1=(2,10),$

$A_2=(2,5),$

$A_3=(8,4),$

$A_4=(5,8),$

$A_5=(7,5),$

$A_6=(6,4),$

$A_7=(1,2),$

$A_8=(4,9)$

k-means

Solving the math..

Step 1: Suppose that the initial seeds (centers of each cluster) are A1, A4 and A7.

i.e. **seed1=A1=(2,10), seed2=A4=(5,8), seed3=A7=(1,2)**

Step 2: Find the Euclidian distance between each seed and point.

$$d(a,b)=\text{sqrt}((x_b-x_a)^2+(y_b-y_a)^2)$$

k-means

Solving the math..

A1:

$$d(A1, \text{seed1})=0 \text{ as } A1 \text{ is seed1}$$

$$d(A1, \text{seed2})= \sqrt{13} >0$$

$$d(A1, \text{seed3})= \sqrt{65} >0$$

→ A1 ∈ cluster1

A3:

$$d(A3, \text{seed1})= \sqrt{36} = 6$$

$$d(A3, \text{seed2})= \sqrt{25} = 5 \quad \leftarrow \text{smaller}$$

$$d(A3, \text{seed3})= \sqrt{53} = 7.28$$

→ A3 ∈ cluster2

A5:

$$d(A5, \text{seed1})= \sqrt{50} = 7.07$$

A2:

$$d(A2, \text{seed1})= \sqrt{25} = 5$$

$$d(A2, \text{seed2})= \sqrt{18} = 4.24$$

$$d(A2, \text{seed3})= \sqrt{10} = 3.16 \quad \leftarrow \text{smaller}$$

→ A2 ∈ cluster3

A4:

$$d(A4, \text{seed1})= \sqrt{13}$$

$$d(A4, \text{seed2})=0 \text{ as } A4 \text{ is seed2}$$

$$d(A4, \text{seed3})= \sqrt{52} >0$$

→ A4 ∈ cluster2

A6:

$$d(A6, \text{seed1})= \sqrt{52} = 7.21$$

k-means

Solving the math..

$$d(A5, \text{seed2}) = \sqrt{13} = 3.60 \leftarrow \text{smaller}$$

$$d(A5, \text{seed3}) = \sqrt{45} = 6.70$$

→ $A5 \in \text{cluster2}$

$$d(A6, \text{seed2}) = \sqrt{17} = 4.12 \leftarrow \text{smaller}$$

$$d(A6, \text{seed3}) = \sqrt{29} = 5.38$$

→ $A6 \in \text{cluster2}$

A7:

$$d(A7, \text{seed1}) = \sqrt{65} > 0$$

$$d(A7, \text{seed2}) = \sqrt{52} > 0$$

$$d(A7, \text{seed3}) = 0 \text{ as } A7 \text{ is seed3}$$

→ $A7 \in \text{cluster3}$

A8:

$$d(A8, \text{seed1}) = \sqrt{5}$$

$$d(A8, \text{seed2}) = \sqrt{2} \leftarrow \text{smaller}$$

$$d(A8, \text{seed3}) = \sqrt{58}$$

→ $A8 \in \text{cluster2}$

end of epoch1

new clusters: 1: $\{A1\}$, 2: $\{A3, A4, A5, A6, A8\}$, 3: $\{A2, A7\}$

k-means

Solving the math..

Step 3: Find the centers of the new clusters: $C1 = (2, 10)$,

$C2 = ((8+5+7+6+4)/5, (4+8+5+4+9)/5) = (6, 6)$,

$C3 = ((2+1)/2, (5+2)/2) = (1.5, 3.5)$

Step 4: After the 2nd epoch the results would be:

1: {A1, A8}, 2: {A3, A4, A5, A6}, 3: {A2, A7}

with centers $C1 = (3, 9.5)$, $C2 = (6.5, 5.25)$ and $C3 = (1.5, 3.5)$.




After the 3rd epoch, the results would be:

1: {A1, A4, A8}, 2: {A3, A5, A6}, 3: {A2, A7}

with centers $C1 = (3.66, 9)$, $C2 = (7, 4.33)$ and $C3 = (1.5, 3.5)$.

k-means

Solving the math..



k-means

Why it's worthy?


- **Simple and Fast:** K-means is easy to understand and implement. It's also computationally efficient for large datasets.
- **Scalable:** K-means works well even when the dataset grows large in size or number of features.
- **Versatile:** K-means is widely applicable to different types of data (e.g., customer segmentation, image compression, text analysis).

k-means

Challenges?

Choosing K:

- Choosing the right number of clusters (K) can be tricky. A common technique is the Elbow Method, where you plot the sum of squared errors (SSE) for different values of K and look for the "elbow point," where increasing K no longer significantly improves the fit.



Cluster Initialization:

- K-means is sensitive to the initial position of centroids. A bad initialization can lead to poor clustering. To solve this, we use K-means++ initialization, which spreads out the initial centroids more intelligently.

Future of AI



আপনার সন্তানকে
মেশিন লার্নিং শিক্ষা দিন

SUFIAN SIDDIQUI